

# Compiler Flags for Code Size and Performance

ADITYA KUMAR



@\_hiraditya\_

# Code size impact on download of mobile apps [2022]

---



Energy cost of download\*:  
0.072 kWh/GB



Total apps downloaded in  
2022\*\*: 143.0B



Energy savings for 1MB  
reduction on each download:  
 **$1\text{MB} * 143\text{B} * 0.072\text{kWh/GB}$**   
**= 10.3 gWh**



Equivalent to 765 million tree  
seedlings grown for 10  
years\*\*\*

*TLDR: Code size matters...*

\* <https://www.carbonbrief.org/factcheck-what-is-the-carbon-footprint-of-streaming-video-on-netflix/>

\*\* <https://www.businessofapps.com/data/app-statistics/>

\*\*\* <https://www.epa.gov/energy/greenhouse-gas-equivalencies-calculator#results>

# Measurement techniques for code size

---

size (binutils)

strings (binutils)

bloaty (<https://github.com/google/bloaty>)

# Measurement techniques

## Size

```
•size gcc/11/libstdc++.dylib
•  __TEXT    __DATA    __OBJC    others    dec    hex
•1703936    65536     0      1851392    3620864    374000
```

## Strings

```
•strings gcc/11/libstdc++.dylib
•2180 strings totalling 36kb
```

## Bloaty

```
•bloaty gcc/11/libstdc++.dylib
•  FILE SIZE      VM SIZE
•-----
• 29.1% 1.00Mi 29.0% 1.00Mi  __TEXT,__text
• 25.0% 882Ki 25.0% 882Ki  String Table
• 16.6% 583Ki 16.5% 583Ki  Symbol Table
• 12.3% 433Ki 12.2% 433Ki  __TEXT,__eh_frame
• 5.0% 176Ki 5.0% 176Ki  Export Info
• 4.1% 146Ki 4.1% 146Ki  __TEXT,__const
• 2.5% 87.8Ki 2.5% 87.8Ki  Weak Binding Info
• 1.2% 41.6Ki 1.2% 41.6Ki  __DATA,__gcc_except_tab
• 1.0% 36.9Ki 1.0% 36.9Ki  __DATA_CONST,__const
• 0.9% 33.3Ki 0.9% 33.3Ki  __TEXT,__text_cold
• 0.5% 16.1Ki 0.5% 16.1Ki  [10 Others]
• 0.5% 15.9Ki 0.0% 945  [__DATA]
• 0.4% 15.0Ki 0.4% 15.0Ki  __TEXT,__cstring
• 0.0% 4 0.3% 11.3Ki  [__LINKEDIT]
• 0.0% 0 0.2% 8.12Ki  __DATA,__bss
• 0.2% 8.01Ki 0.2% 8.01Ki  [__DATA_CONST]
• 0.2% 7.43Ki 0.2% 7.43Ki  Function Start Addresses
• 0.0% 0 0.2% 6.88Ki  __DATA,__common
• 0.2% 6.08Ki 0.2% 6.08Ki  Indirect Symbol Table
• 0.1% 4.59Ki 0.1% 4.59Ki  __DATA,__la_symbol_ptr
• 0.1% 3.44Ki 0.1% 3.44Ki  __TEXT,__stubs
• 100.0% 3.44Mi 100.0% 3.45Mi  TOTAL
```

# Code size optimization flags

---

## Always decrease code size

- `-Os`
- `-Oz` (<https://godbolt.org/z/zdMKhfe7G>)
- `-flto`
- `-fno-unroll-loops`
- `-fno-exceptions`
- `-fno-rtti`

## May increase code size in some codebase

- `-function-sections -Wl,--gc-sections`
- `-fno-jump-tables`

# Additional Code size optimization flags

thin-lto (-flto=thin llvm)

GVNHoist (-mllvm --enable-gvnhoist)

GVNSink (-mllvm --enable-gvnsink)

Machine outliner (-mllvm -enable-machine-outliner)

Hot cold splitting (-mllvm -hot-cold-split)

Change Inliner threshold

- -mllvm -inline-threshold=n

# Impact of performance optimizations

---

- On smartphones
  - Android, iPhone
- Power consumption (4kWh per year)
  - Very small number but..
  - \* 6.8 billion
- Improving 1% of power consumption on all smartphones
  - $1\% * 4\text{kWh} * 6.8 \text{ billion} = 272 \text{ gWh} = 20 \text{ billion tree seedlings grown for 10 years}$

# Compiler optimizations for performance

---

Optimization level -O2, -O3, -Ofast (<https://godbolt.org/z/4vPr6dhWb>)

LTO (-flto, -flto=thin)

-fno-exceptions

-fno-rtti

-mllvm -hot-cold-split (<https://godbolt.org/z/zhddPhSex>)

-mllvm --enable-gvnhoist

-mllvm --enable-gvnsink



# Performance analysis tools

- Valgrind
- Linux Perf
- Visual studio performance tool
- Intel Vtune
- XCode Instruments

# Performance Analysis with Valgrind

---

`valgrind [--tool=memcheck]`

- valgrind mostly known for its memory leak checker

`valgrind --tool=cachegrind`

- cache and branch simulator
- count read, write, and branch instructions

`valgrind --tool=callgrind`

- execution call graph
- visualization tool kcachegrind

# Valgrind: Example – SQLite

```
$ valgrind --tool=cachegrind ./sqlite_llvm <test.sql >/dev/null  
[...]
```

Ir	I1mr	ILmr	Dr	D1mr	DLmr	Dw	D1mw	DLmw	
1,278,771,731	29,231,219	35,783	359,414,267	6,707,514	528,920	197,515,528	2,594,262	171,968	PROGRAM TOTALS

Ir	I1mr	ILmr	Dr	D1mr	DLmr	Dw	D1mw	DLmw	file:function
363,052,233	7,560,087	3,122	97,707,865	1,084,529	77,197	44,505,055	217,826	29,838	src/sqlite3.c:sqlite3VdbeExec
95,048,357	80,721	111	33,248,107	59,086	7,273	20,173,275	91	7	src/sqlite3.c:vdbeRecordCompareWithSkip
68,045,026	695,509	1,144	14,883,933	114,698	1,918	5,525,733	272,507	19,249	src/sqlite3.c:balance
56,713,554	1,101,002	276	18,416,705	683,914	21,085	3,453,665	1,947	25	src/sqlite3.c:sqlite3BtreeMovetoUnpacked
45,344,891	59,660	66	13,589,490	66,121	18,775	12,795,281	59,451	86	src/sqlite3.c:sqlite3VdbeRecordUnpack
36,550,248	47,192	94	9,615,816	217,845	11,567	0	0	0	src/sqlite3.c:cellSizePtr
35,156,491	1,031,905	859	7,810,853	489,509	1,936	6,546,085	175,469	26,159	/build/glibc-2.19/malloc/malloc.c:_int_malloc
34,402,967	219,015	40	12,316,213	31,625	1,007	0	0	0	src/sqlite3.c:vdbeRecordCompareInt

# Performance Analysis with Linux Perf

---

## perf stat

- sum up all counters

## perf record

- record events

## perf report

- Shows the profile collected using `perf record`

# Perf stat: Example – SQLite

---

```
$ perf stat ./sqlite_llvm <test.sql >/dev/null
```

```
Performance counter stats for './sqlite_llvm':
```

1045.856070	task-clock (msec)	#	1.000 CPUs utilized	
1	context-switches	#	0.001 K/sec	
0	cpu-migrations	#	0.000 K/sec	
809	page-faults	#	0.774 K/sec	
1,636,720,010	cycles	#	1.565 GHz	[83.16%]
548,530,227	stalled-cycles-frontend	#	33.51% frontend cycles idle	[83.16%]
218,991,051	stalled-cycles-backend	#	13.38% backend cycles idle	[67.04%]
3,385,841,295	instructions	#	2.07 insns per cycle	
		#	0.16 stalled cycles per insn	[83.54%]
709,436,490	branches	#	678.331 M/sec	[83.54%]
2,586,354	branch-misses	#	0.36% of all branches	[83.17%]

```
1.045918998 seconds time elapsed
```

# Perf record: Example – xalancbmk

---

```
$ perf record ./xalancbmk
$ perf report
0.20 629a84:  ldr    w9, [x0,#24]
18.71 629a88:  ldr    w8, [x1,#24]
12.93 629a8c:  cmp    w9, w8
2.74 629a90:  b.ne   629af8 <xalanc_1_8::XalanDOMString::equals
2.00 629a94:  ldp    x8, x10, [x0]
2.43 629a98:  cmp    x8, x10
1.80 629a9c:  ldp    x10, x12, [x1]
1.03 629aa0:  adrp   x11, 704000 <vtable for xalanc_1_8::ReusableArenaBlock+0x8>
0.53 629aa4:  add    x11, x11, #0xb08
0.03 629aa8:  csel   x8, x11, x8, eq
1.33 629aac:  cmp    x10, x12
0.34 629ab0:  csel   x10, x11, x10, eq
1.78 629ab4:  cbz    w9, 629b00 <xalanc_1_8::XalanDOMString::equals
0.02 629ab8:  ldrrh  w11, [x8]
4.02 629abc:  ldrrh  w12, [x10]
3.75 629ac0:  cmp    w11, w12
1.03 629ac4:  b.ne   629b08 <xalanc_1_8::XalanDOMString::equals
1.16 629ac8:  lsl    x9, x9, #1
```

# Analyzing System Performance

---

## Vary one Component of the System at a time

- Measure impact of one component on the System
- Run multiple times

## Disable frequency scaling

- cpufrequtils

## Performance metrics

- Dynamic profiles, compiler logs

## Systematic performance analysis

- Monitor performance regression over time
- Time series: track performance of system over time
- Git bisect performance changes

# Performance analysis pitfalls

---

## Central tendencies

- The median instead of the mean

## Use the quantile values instead of a single median value

- Helps with prioritization

## Outlier detection

- Filter outliers

## Weighted samples for combining historical data

- Recent data more important than the previous data.



# Case studies

# std::vector copy constructor

```
// clang++ -std=c++17 -O3 -fno-exceptions -stdlib=libc++
#include<vector>
```

```
using T = int;
T vec_copy(const std::vector<T> &v1) {
    auto v(v1);
    return 10;
}
```

```
/*
// Or
T vec_copy(std::vector<T> v1) {
    auto v(v1);
    return 10;
}
*/
```

```
vec_copy(std::__1::vector<int, std::__1::allocator<int>
> const&): # @vec_copy(std::__1::vector<int,
std::__1::allocator<int> > const&)
push r15
push r14
push rbx
mov r15, qword ptr [rdi]
mov r14, qword ptr [rdi + 8]
mov rbx, r14
sub rbx, r15
je .LBB0_10
js .LBB0_11
mov rdi, rbx
call operator new(unsigned long)@PLT
add rbx, -4
cmp rbx, 28
jb .LBB0_3
mov rcx, rax
sub rcx, r15
cmp rcx, 32
jb .LBB0_3
shr rbx, 2
inc rbx
mov rsi, rbx
and rsi, -8
lea rcx, [r15 + 4*rsi]
lea rdx, [rax + 4*rsi]
xor edi, edi
.LBB0_6: # =>This Inner Loop Header: Depth=1
movups xmm0, xmmword ptr [r15 + 4*rdi]
movups xmm1, xmmword ptr [r15 + 4*rdi + 16]
movups xmmword ptr [rax + 4*rdi], xmm0
movups xmmword ptr [rax + 4*rdi + 16], xmm1
add rdi, 8
cmp rsi, rdi
jne .LBB0_6
cmp rbx, rsi
jne .LBB0_8
jmp .LBB0_9
.LBB0_3:
mov rcx, r15
mov rdx, rax
.LBB0_8: # =>This Inner Loop Header: Depth=1
mov esi, dword ptr [rcx]
mov dword ptr [rdx], esi
add rcx, 4
add rdx, 4
cmp rcx, r14
jne .LBB0_8
.LBB0_9:
mov rdi, rax
call operator delete(void*)@PLT
.LBB0_10:
mov rax, 10
pop rbp
```

# [copy] After optimization D147741

---

```
// clang++ -std=c++17 -O3 -fno-  
exceptions -stdlib=libc++  
#include<vector>  
  
using T = int;  
T vec_copy(const std::vector<T> &v1)  
{  
    auto v(v1);  
    return 10;  
}
```

<https://godbolt.org/z/7q9PdrT81>

```
vec_copy(std::__1::vector<int,  
std::__1::allocator<int> > const&):  
sub rsp, 24  
xorps xmm0, xmm0  
movaps xmmword ptr [rsp], xmm0  
mov qword ptr [rsp + 16], 0  
mov rax, qword ptr [rdi + 8]  
sub rax, qword ptr [rdi]  
je .LBB0_2  
js .LBB0_3
```

```
.LBB0_2:  
mov eax, 10  
add rsp, 24  
ret
```

```
.LBB0_3:  
mov rdi, rsp  
call std::__1::vector<int,  
std::__1::allocator<int>  
>::__throw_length_error[abi:v170000]()  
const  
.L.str:  
.asciz "vector"
```

```
.L.str.1:  
.asciz "length_error was thrown in -fno-  
exceptions mode with message \"%s\""
```

# std::vector<int> access

---

```
#include <vector>
void f(int);

void subscript_operator(std::vector<int> v)
{
    for (std::vector<int>::size_type i = 0; i < v.size(); i++)
        f(v[i]);
}

void iterator(std::vector<int> v)
{
    for (std::vector<int>::const_iterator i = v.begin(); i != v.end(); ++i)
        f(*i);
}

void range_loop(std::vector<int> v)
{
    for (auto i : v)
        f(i);
}
```

# std::vector<int> access

```
#include <vector>
void f(int);
```

```
void subscript_operator(std::vector<int> v)
{
    for (std::vector<int>::size_type i = 0; i < v.size(); i++)
        f(v[i]);
}
```

```
void iterator(std::vector<int> v)
{
    for (std::vector<int>::const_iterator i = v.begin(); i != v.end(); ++i)
        f(*i);
}
```

```
void range_loop(std::vector<int> v)
{
    for (auto i : v)
        f(i);
}
```

```
.LBB1_1: # =>Inner Loop Header
mov edi, dword ptr [rbx]
call f(int)@PLT
add rbx, 4
cmp rbx, r14
jne .LBB1_1
```

```
.LBB0_2: # => Inner Loop Header:
mov edi, dword ptr [rax + 4*r14]
call f(int)@PLT
inc r14
mov rax, qword ptr [rbx]
mov rcx, qword ptr [rbx + 8]
sub rcx, rax
sar rcx, 2
cmp r14, rcx
jb .LBB0_2
```

```
.LBB2_2: # =>Inner Loop Header:
mov edi, dword ptr [r14]
call f(int)@PLT
add r14, 4
cmp r14, qword ptr [rbx + 8]
jne .LBB2_2
```

<https://godbolt.org/z/zMGTEfYM>

# References (performance)

---

- <https://github.com/hiraditya/std-benchmark/blob/master/docs/slides/slide-DAC-2017.pdf>
- [Profile-based Indirect Call Promotion](#)
- [Developer Tools #WWDC16](#)
- [Eliminating Virtual Function Calls in C++ Programs](#)
- [http://boost-sandbox.sourceforge.net/libs/proto/doc/html/boost\\_proto/users\\_guide/getting\\_started/hello\\_world.html](http://boost-sandbox.sourceforge.net/libs/proto/doc/html/boost_proto/users_guide/getting_started/hello_world.html)
- [Linkers and Loaders by John R. Levine](#)
- <https://man7.org/linux/man-pages/man2/getrusage.2.html>
- <https://www.gnu.org/software/hurd/gnumach-doc/Task-Information.html>
- [https://en.wikipedia.org/wiki/Lazy\\_initialization](https://en.wikipedia.org/wiki/Lazy_initialization)
- <https://stackoverflow.com/a/28146199/811335>
- <http://www.dbp-consulting.com/tutorials/debugging/linuxProgramStartup.html>
- <https://gist.github.com/C0deH4cker/80b53de22012146ea9d8>
- <https://blog.timac.org/2016/0804-dump-decrypted-mach-o-files/>
- <https://embeddedartistry.com/blog/2019/05/20/exploring-startup-implementations-os-x/>
- <https://engineering.fb.com/2015/11/20/ios/optimizing-facebook-for-ios-start-time/>

# References (code size)

---

- <https://github.com/hiraditya/std-benchmark/blob/master/docs/slides/CppConCodesizeCompilerOptimizationAndTechniques.pdf>
- `man gcc`
- `clang --help-hidden`
- "-Os Matters" by Mark Zeren
  - <https://www.youtube.com/watch?v=vGV5u1nxqd8>
- <https://github.com/google/bloaty>
- <https://www.mail-archive.com/gcc@gcc.gnu.org/msg91116.html>
- <http://gcc.gnu.org/>

# Compiler Flags for Code Size and Performance

ADITYA KUMAR



@\_hiraditya\_