


Compiler optimizations and software engineering techniques for improving application launch time

Aditya Kumar
Software Engineer

 @_hiraditya_
[linkedin.com/in/hiraditya](https://www.linkedin.com/in/hiraditya)

Program startup types

- Cold start
- Warm start
 - Cached/Prewarmed state
- Hot start

App startup compiler optimizations and techniques [cold start]

Program launch sequence

- Loader
- Rebase+Binding
- Static initializer
- crt etc.

Factors affecting startup


- Page faults
- extern functions
- Language runtime
- Static initializers and attribute constructors

Measurement and optimization techniques of different stages of startup

- Build time and runtime measurement techniques
- Build time and runtime optimizations

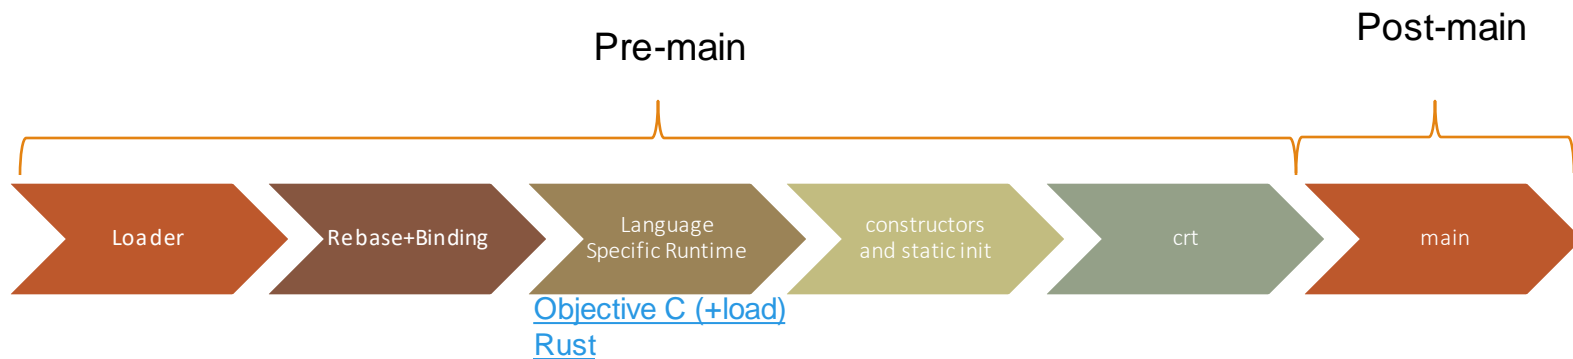
Compiler optimizations and software engineering techniques

- Compiler optimizations that helps with startup
- Software engineering techniques
- Compiler optimizations yet to be implemented in LLVM



Program launch sequence

Program launch sequence



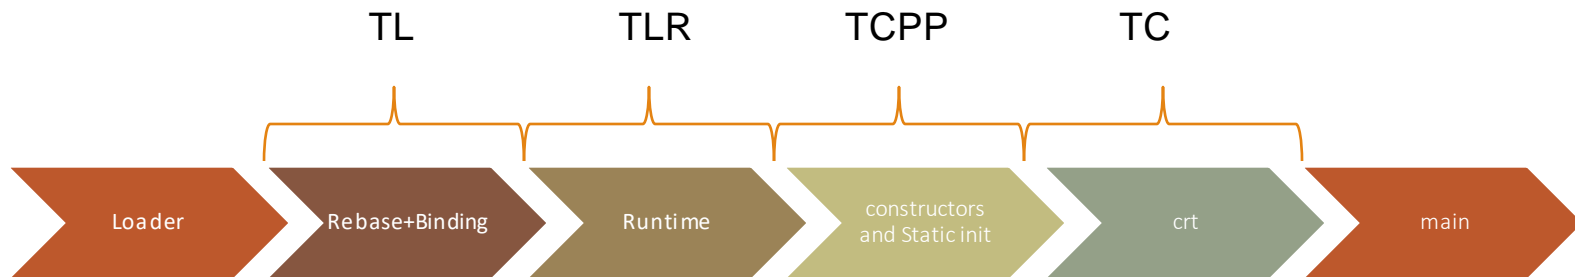
Measuring launch time spans

- By inserting timers wherever possible

Helps with

- Detect regressions
- Find bottlenecks

Measuring launch time spans



TL = Time to finish Loading the binary into memory

TLR = Time taken to execute Language runtime code

TCPP = Time taken by C++ static initializers and static constructors

TC = Time taken by the C runtime



Factors affecting cold start

Launch time is affected by

- TL
 - Page faults
 - extern functions
- TLR
 - Language runtime, pre-main methods
- TCPP
 - Static initializers and attribute constructors
 - C++ standard library static constructors
- TC
 - C runtime library setup


TL = Time to finish Loading the binary into memory
TLR = Time taken to execute Language runtime code
TCPP = Time taken by C++ static initializers and static constructors
TC = Time taken by the C runtime

Measurement & optimization techniques

Launch time can be improved by

- TL
 - Reducing Page faults
 - Reducing number of rebases and bindings
- TLR
 - Reducing work done by the language runtime
- TCPP
 - Reducing number of static initializers and static constructors
- TC
 - N/A for the most part

TL = Time to finish Loading the binary into memory
TLR = Time taken to execute Language runtime code
TCPP = Time taken by C++ static initializers and static constructors
TC = Time taken by the C runtime



[TL] Page faults, rebases and bindings

Measuring page faults

- `getrusage`
 - `struct rusage::ru_majflt`
- `task_info`
 - `struct task_events_info::faults`
- `ps`
 - `majflt`

Compiler optimizations that helps reduce page faults

- Compiler optimizations to reduce binary size (page faults \propto binary size):
 - `-Oz`, `-Os`
 - `-fno-exceptions`
 - `-fno-rtti`
 - `-flto`, `-flto=thin`
 - Identical code folding (function merging)
 - Machine outliner
 - Devirtualization
- Reduce working set size
 - hot-cold-splitting

Other
optimizations
that helps
reduce page
faults

- Order file
- Code restructuring and annotations

Order file

- Helps with spatial locality
- Order file is a line separated list of function names
 - Mangled names for C++

```
$ clang++ hello.cpp -o a.out -O2
$ objdump -D a.out | bash_magic
__main
__put_character_sequence
__pad_and_output
```

```
$ cat hello.order
__ZNSt3__124__put_character_sequenceIc...
__ZNSt3__116__pad_and_outputIcNS_11char_traitsI...
__main
```

```
$ clang++ hello.cpp -o b.out -O2 -Wl,-order_file -Wl,hello.order
$ objdump -D b.out | bash_magic
__put_character_sequence
__pad_and_output
__main
```


Steps to generate order file

- Execute the program
 - Instrumented or Bare
- Collect traces
 - Sampling (ftrace, dtrace)
 - Instrumented (-finstrument-function-entry-bare, -forder-file-instrumentation)
 - lldb (set breakpoint on all the functions)
- Generate order file by merging traces
 - Functions, static initializers
- Build the app with order file
 - `ld -order_file <file> -order_file_statistics <args> <binaries> -o a.out`

Measuring rebases and bindings

- Rebases
 - `objdump [-macho] --rebase Binary | wc`
- Bindings
 - `llvm-objdump [-macho] --bind Binary | wc`
 - `llvm-objdump [-macho] --lazy-bind Binary | wc`
 - `llvm-objdump [-macho] --weak-bind Binary | wc`
- Using dyldinfo
 - `dyldinfo [-rebase|-bind|-weak-bind|...] | wc`

Reducing number of rebases and bindings

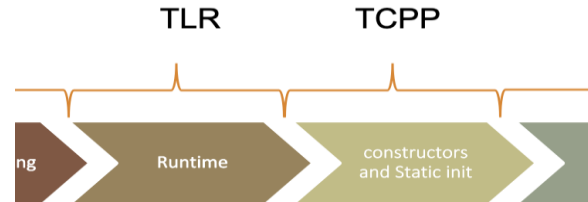
- Compiler optimizations to reduce rebases and bindings:
 - `-flto`, `-flto=thin`
- Making extern functions static when possible

The background of the slide features a pattern of interlocking puzzle pieces. One piece in the center is a solid dark grey, while the others are light grey with thin white outlines. A thin horizontal line is positioned above the text, and a solid orange horizontal bar is at the bottom of the slide.

[TLR] Language Runtime

Measuring impact of language runtime

- Use debugger or binutil tools to find the first called function
- Add 'start' timer
 - In the first function called.
- The 'end' timer
 - In the last function called or,
 - In the first function of static initializer (next stage)



Reducing language runtime

- Removing pre-main functions as much as possible
 - Objective-C: Reduce number of +load in your program
- Do minimal work in pre-main functions
- Also helps with reliability of the app

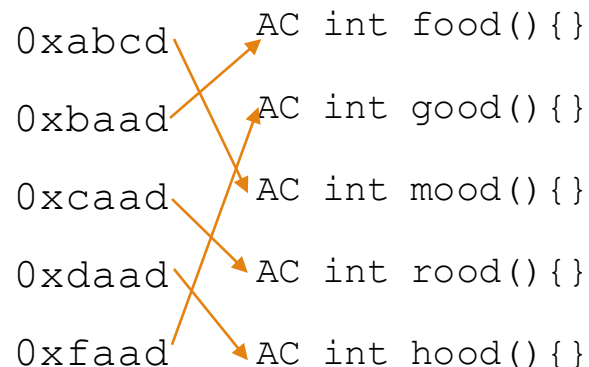


[TCPP] Static initializers, constructors



Init array & SIOF

- Pointers to functions stored as array
 - Static constructor
- Each function in the init array is invoked in a sequence
- The order of init array is non-deterministic
 - Depends on the order of definition, as well as order of linking



Measuring impact of static initializers

- Approach #1:
 - Find which function executes first, and last (with debugger, or binutil tools)
 - Insert timers
- Approach #2
 - Add two constructors (begin, end) with init_priority
 - [gnu linker] respects priority across translation units
 - Manipulate the init_array (mod_init_func) to reorder the constructors to begin and end of the array
 - Good Starting point:
<https://gist.github.com/C0deH4cker/80b53de22012146ea9d8>

Reducing number of static initializers and attribute constructors

- Removing global static objects
 - Lazy evaluation by initializing the static object in function [\[Meyer's singleton\]](#)
- Removing unnecessary
`__attribute__((constructors))`

The background of the slide features a light gray puzzle piece pattern. A large, semi-transparent puzzle piece is centered behind the text. Overlaid on this are thin, light gray wavy lines that create a sense of motion or flow. The title text is centered and framed by two horizontal lines.

Compiler Optimizations

Compiler optimizations that affect startup

- Optimization level -Os, -Oz, -O3
- LTO
- -fno-exceptions
- -fno-rtti
- -mllvm -hot-cold-split=true
- -fmerge-functions

Hot cold splitting (-mllvm -hot-cold-split=true)

- Outlines cold subgraphs and puts them in a separate section
- Reduces page fault
- Enabled for all iPhone apps (in swift-llvm)

Profile guided indirect call promotion

- Helps with inlining and other optimization
- Enabled in lto and thinlto

Devirtualization

- Whole program devirtualization
 - `-flto, -flto=thin`
- [Objective-C] Annotating classes+methods with `objc_direct`.
 - Calls to `objc_direct` methods become C function calls
 - Objective-C metadata w.r.t. methods are removed



Software Engg Techniques

Code restructuring and annotations

Early evaluation

- `constexpr` (C++)

Lazy evaluation

- Creating functions objects
- Set of functions -> `HashMap<Key, Lambda Function>`
- Moving the code not used in startup to a separate shared library

Caching


- Possibly with value profiling

Codegen at runtime


- `vector<lambda>`
- `unordered_map<key, lambda>`

Removing Dead Code

- Remove test code shipped to production
 - Code used only in tests ending up in the shipped binary
 - `nm <Binary> | grep -i "test\|debug" :)`



Compiler Optimizations yet to be implemented

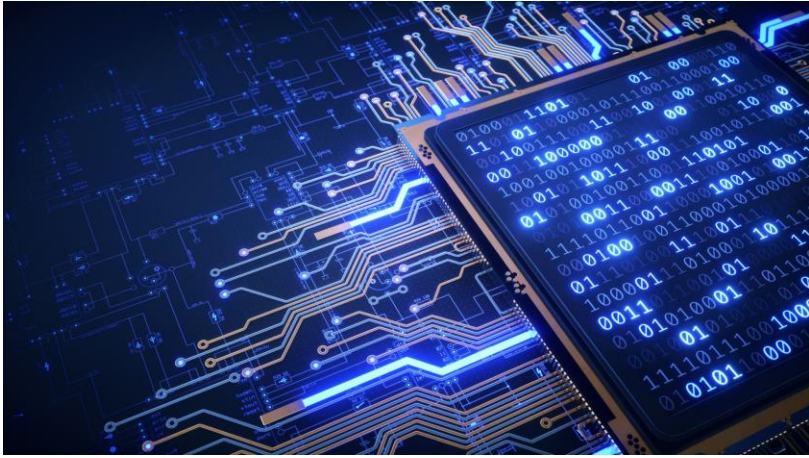


Measuring static init time

- Limitations of linker, SIOF, C++
- Binary reordering of static initialization sequence


References

- [Profile-based Indirect Call Promotion](#)
- [Developer Tools #WWDC16](#)
- [Eliminating Virtual Function Calls in C++ Programs](#)
- http://boost-sandbox.sourceforge.net/libs/proto/doc/html/boost_proto/users_guide/getting_started/hello_world.html
- [Linkers and Loaders by John R. Levine](#)
- <https://man7.org/linux/man-pages/man2/getrusage.2.html>
- <https://www.gnu.org/software/hurd/gnumach-doc/Task-Information.html>
- https://en.wikipedia.org/wiki/Lazy_initialization
- <https://stackoverflow.com/a/28146199/811335>
- <http://www.dbp-consulting.com/tutorials/debugging/linuxProgramStartup.html>
- <https://gist.github.com/C0deH4cker/80b53de22012146ea9d8>
- <https://blog.timac.org/2016/0804-dump-decrypt-mach-o-files/>
- <https://embeddedartistry.com/blog/2019/05/20/exploring-startup-implementations-os-x/>
- <https://engineering.fb.com/2015/11/20/ios/optimizing-facebook-for-ios-start-time/>



App Startup compiler optimizations and techniques for embedded systems

Aditya Kumar
Software Engineer

 @_hiraditya_
[linkedin.com/in/hiraditya](https://www.linkedin.com/in/hiraditya)

Backup
